ARYAN MEHTA

Machine Learning Engineer

aryan.mehta@email.com | +91 98765 43210 | linkedin.com/in/aryanmehta | github.com/aryan-ml | Gurgaon, Haryana

PROFESSIONAL SUMMARY

Results-driven Machine Learning Engineer with 3+ years of experience designing, building, and deploying production-grade ML systems. Strong expertise in NLP, deep learning, and scalable MLOps pipelines. Passionate about applying AI to real-world problems at scale. Experienced in Python, PyTorch, and cloud platforms (Azure). Seeking to contribute to Microsoft's AI-first mission and build intelligent systems that reach millions of users.

WORK EXPERIENCE

Machine Learning Engineer | Paytm (One97 Communications) *Noida, UP • Jul 2022 – Present*

– Built a real-time fraud detection system using gradient-boosted trees (XGBoost + LightGBM) achieving 97.4% precision, processing 2M+ daily transactions.
– Developed an NLP-based customer intent classifier (BERT fine-tuned on domain data) reducing ticket misrouting by 38%.
– Designed and owned end-to-end MLOps pipeline on Azure ML: automated retraining, model versioning, A/B testing, and performance monitoring.
– Reduced model inference latency by 60% via ONNX conversion and TensorRT optimisation on deployed endpoints.
– Mentored 2 junior engineers; led weekly ML reading groups to keep the team current on research.

ML Intern | Samsung Research Institute – Bangalore *Bengaluru, KA • Jan 2022 – Jun 2022*

– Fine-tuned a multilingual speech recognition model (Conformer-CTC) for 5 Indian languages, achieving 8.2% WER on internal benchmarks.
– Built a data augmentation pipeline (SpecAugment + noise injection) that improved robustness by 12% on noisy environments.
– Collaborated with the on-device AI team to quantise models for Exynos chipsets, reducing model size by 45% with <1% accuracy drop.

Research Intern | IIT Delhi – CSSL Lab *New Delhi • May 2021 – Jul 2021*

– Researched continual learning techniques (EWC, Progressive Nets) for image classification under distribution shift.
– Published findings as co-author in a workshop paper at NeurIPS 2021 (non-archival).

EDUCATION

B.Tech – Computer Science & Engineering | VIT University *Vellore, TN • 2018 – 2022*

– CGPA: 8.9 / 10 | Specialisation: AI & Machine Learning
– Relevant coursework: Deep Learning, NLP, Computer Vision, Algorithms & Data Structures, Probability & Statistics

TECHNICAL SKILLS

Languages	Python (expert), C++, SQL, Bash

ML / DL	PyTorch, TensorFlow, scikit-learn, HuggingFace Transformers, XGBoost, LightGBM
NLP / GenAI	BERT, GPT-2/3, LLaMA, LangChain, RAG pipelines, vector databases (FAISS, Pinecone)
MLOps	Azure ML, MLflow, DVC, Docker, Kubernetes, CI/CD (GitHub Actions), Prometheus
Data / Cloud	Azure (ADF, Synapse, AKS), PySpark, Pandas, NumPy, PostgreSQL, Redis

KEY PROJECTS

DocBot – RAG-based Enterprise Q&A System | *github.com/aryan-ml/docbot*

– Built a production RAG pipeline (LangChain + FAISS + LLaMA-3) that answers natural-language queries over 50K+ internal documents with <2s latency.
– Implemented hybrid retrieval (dense + sparse BM25) improving answer relevancy by 22% over vanilla vector search.

Real-time Anomaly Detection for IoT Streams

– Designed a streaming anomaly detector (Isolation Forest + LSTM autoencoder) on Azure Event Hubs achieving 99.1% recall on manufacturing sensor data.
– Deployed as a containerised microservice on AKS, auto-scaling to handle 500K events/sec with p99 latency of 45ms.

CERTIFICATIONS & PUBLICATIONS

– Microsoft Certified: Azure AI Engineer Associate (AI-102) – 2023
– Deep Learning Specialisation – Andrew Ng, Coursera – 2021
– Co-author: 'Mitigating Catastrophic Forgetting via Dynamic Sparse Masks' – NeurIPS 2021 Workshop (DistributeML)

ACHIEVEMENTS

– Ranked Top 50 globally in Kaggle's "CommonLit Readability" NLP competition (out of 3,633 teams)
– Won 1st place – Smart India Hackathon 2021 (AI Track) for an AI-powered crop disease detection app
– Open-source: DocBot has 1,200+ GitHub stars and 140+ forks